

Comprehensive Analysis of Semi-Supervised Self-Training Architectures (2022–2026)

The Paradigm Shift in Semi-Supervised Learning and Self-Training

Semi-supervised learning has historically operated on a fundamental mathematical premise: leveraging a restricted corpus of labeled data alongside a vast, unstructured repository of unlabeled data to define highly robust decision boundaries in low-density regions. Central to this domain is self-training, a process whereby a machine learning model generates pseudo-labels for unlabeled instances and iteratively optimizes its parameters upon its own predictions. Over the past four years, as extensively chronicled by the proceedings of the Conference on Neural Information Processing Systems (NeurIPS), the International Conference on Machine Learning (ICML), and the International Conference on Learning Representations (ICLR), the self-training paradigm has undergone a profound architectural and theoretical transformation.¹

Historically, self-training algorithms relied heavily on static heuristics. Models utilized fixed confidence thresholds to filter pseudo-labels, operating under the naive assumption that high-confidence predictions correlate intrinsically with high accuracy. This approach, while effective in controlled environments with balanced class distributions, degrades catastrophically under conditions of distribution shift, heteroscedastic noise, out-of-distribution unlabelled data, and complex modality architectures.⁴ The research published between 2022 and 2026 demonstrates a systematic departure from these rigid heuristics toward dynamic, adaptive, and theoretically bounded frameworks.⁷

The field has expanded far beyond the classical boundaries of standard image classification. Contemporary research has vigorously pushed self-training into semi-supervised continuous regression, dense pixel-level segmentation, cross-domain adaptation, graph consistency learning, and the iterative, autonomous self-improvement of large language models and vision foundation models.⁸ Concurrently, the theoretical understanding of self-training has deepened. Researchers have moved beyond empirical observation to provide rigorous mathematical proofs explaining why pseudo-labeling alters gradient dynamics, how recursive self-training induces Model Autophagy Disorder (MAD), and why non-stationary training dynamics often yield superior optimization signals compared to stationary convergence.¹²

This exhaustive report provides a systematic analysis of the accepted papers at NeurIPS, ICML, and ICLR from 2022 to 2026. It dissects the transition from basic, confidence-based pseudo-labeling to advanced bi-level optimization techniques. Furthermore, it examines the

integration of parameter-efficient fine-tuning with foundation models, the mechanisms developed to counter long-tailed distribution mismatches, and the emerging theoretical frontiers that address model collapse, negative extrapolation, and self-penalization in reinforcement learning architectures.

Theoretical Foundations and the Mathematics of Pseudo-Labeling

To fully grasp the advancements in self-training methodologies, one must first examine the theoretical bedrock established during this period. For years, the empirical success of pseudo-labeling outpaced its theoretical justification. The ICLR 2023 paper titled "How Does Semi-supervised Learning with Pseudo-labelers Work? A Case Study" provided a landmark unified mathematical proof analyzing the dynamics of pre-training followed by linear probing.¹² The researchers investigated a specific two-stage semi-supervised learning approach: first, a two-layer convolutional neural network is trained on unlabeled data with the assistance of a pseudo-labeler, and subsequently, the pre-trained network is linearly probed using a limited quantity of labeled data.¹²

By operating under a defined toy data generation model where noise patches are sampled from Gaussian distributions, the authors successfully manipulated gradient formulations to establish rigorous bounds on signal growth rates versus noise growth rates.¹⁵ They demonstrated mathematically that, under conditions of sufficient over-parameterization, the semi-supervised approach utilizing pseudo-labels can achieve near-zero test loss.¹² In stark contrast, a neural network trained via purely supervised learning on the exact same limited quantity of labeled data is mathematically restricted to achieving only a constant test loss.¹² This provided the definitive theoretical separation demonstrating that pseudo-labeling fundamentally alters the trajectory of gradient descent, formally validating the immense empirical successes observed across the discipline.

Building upon the theoretical boundaries of empirical risk minimization, research at NeurIPS 2025 expanded into the complexities of covariate shift. The paper "Pseudo-Labeling for Kernel Ridge Regression under Covariate Shift" developed a principled approach to target-distribution regression when labeled source data possesses a significantly different feature distribution.¹⁸ The proposed non-asymptotic excess risk bounds demonstrate that splitting labeled data to construct candidate models and an imputation model allows for optimal adaptation. The authors quantified this adaptation through a novel definition of effective sample size, proving that using pseudo-labels for model selection does not significantly hinder performance while achieving minimax optimal error rates up to a polylogarithmic factor.¹⁸

The theoretical understanding of adversarial robustness within self-training loops was formalized in the ICML 2025 paper "ASTrA: Adversarial Self-supervised Training with Adaptive-Attacks".¹⁹ Adversarial self-supervised training typically requires immense computational overhead and relies heavily on static, pre-defined attack parameters. ASTrA

introduced a learnable, self-supervised attack strategy network that functions autonomously, discovering optimal adversarial attack parameters through a continuous exploration-exploitation process within a single training episode.¹⁹ By theoretically bounding the exploration phase, the framework demonstrated high resilience to robust overfitting while maintaining strong semi-supervised performance, proving that dynamic adversarial perturbations can act as a stabilizing regularizer in self-training pipelines.¹⁹

Overcoming the Quantity-Quality Tradeoff: Continuous and Adaptive Thresholding

The core vulnerability of classical self-training, particularly frameworks like FixMatch, lies in the strict quantity-quality tradeoff. Setting a high confidence threshold ensures that only highly accurate pseudo-labels are incorporated into the training loop; however, it severely limits the quantity of usable data, often entirely ignoring hard-to-learn classes.⁵ Conversely, a low threshold maximizes data utilization but introduces severe confirmation bias, wherein the model cyclically reinforces its own erroneous predictions.⁵

The introduction of FreeMatch at ICLR 2023 marked a pivotal milestone in resolving this bottleneck.⁴ FreeMatch entirely abandoned the reliance on pre-defined thresholds or ad-hoc, manually designed thresholding schedules. Instead, the authors recognized that the optimal threshold must reflect the model's dynamic learning status.⁴ They introduced Self-Adaptive Thresholding (SAT), a mechanism that dynamically computes the confidence threshold based on the moving average of the model's prediction confidences.²⁶ As the model's epistemic uncertainty decreases over the course of training, the threshold adjusts organically. Furthermore, because underlying pseudo-label distributions are rarely uniform, FreeMatch incorporated a Self-Adaptive class Fairness (SAF) regularization metric.²⁶ This metric normalizes the expectation of probability by the histogram distribution of the pseudo-labels, effectively countering the negative effects of class imbalance. The empirical results were profound: FreeMatch achieved massive error rate reductions, including a 5.78% reduction on CIFAR-10 when labeled data was restricted to an extreme setting of a single sample per class.²⁶

While FreeMatch adaptive thresholds were highly successful, the overarching reliance on binary truncation—where a sample is either included or excluded—remained a structural limitation. SoftMatch, also presented at ICLR 2023, proposed a continuous alternative to binary thresholding to further optimize the quantity-quality tradeoff.⁵ Rather than strictly filtering out low-confidence samples, SoftMatch introduced a sample weighting function based on a truncated Gaussian distribution.⁵ This function fits the underlying confidence distribution of the model's historical predictions. It assigns continuous, exponentially decaying lower weights to pseudo-labels according to the deviation of their confidence from the mean of the Gaussian.⁵ By doing so, SoftMatch effectively leverages unconfident yet potentially correct pseudo-labels, ensuring that the model does not discard valuable gradient information regarding complex, weakly-learned classes.²²

The optimization of pseudo-labeling parameters was further abstracted mathematically in the PabLO framework, presented at NeurIPS 2024.⁷ Researchers identified that standard confidence metrics, such as Maximum Softmax Probability (MSP), are inherently flawed because neural networks can produce arbitrarily high confidence scores for out-of-distribution inputs.³¹ PabLO formulated the pseudo-labeling process not as a heuristic design choice, but as a strict optimization problem over a flexible functional space of confidence scoring functions and thresholds.⁷ Expanding upon threshold-based auto-labeling (TBAL) systems, PabLO learns optimal scoring functions explicitly aligned with the downstream semi-supervised task, eliminating the need for manual heuristic designs.⁷ This framework integrates seamlessly into existing pipelines, mathematically bounding the quality of the pseudo-labels, and yielding up to an 11% improvement in test accuracy over standard baselines while requiring substantially fewer training iterations.⁷

Addressing Distribution Mismatches and Long-Tailed Learning

A persistent, unresolved challenge in real-world semi-supervised learning is the phenomenon of distribution mismatch, where the underlying class distribution of the labeled dataset diverges significantly from that of the unlabeled dataset. This issue is particularly acute in long-tailed semi-supervised learning (LTSSL), where head classes naturally dominate the data distribution while tail classes suffer from severe underrepresentation.³⁴

The ICML 2023 paper on "Bidirectional Adaptation for Robust Semi-Supervised Learning with Inconsistent Data Distributions" provided a critical theoretical framework for understanding why standard semi-supervised algorithms fail under these conditions.⁶ The authors mathematically demonstrated that distribution discrepancies between pseudo-label predictions and target predictions lead to severe generalization errors. They identified three specific structural flaws in legacy algorithms: the tight coupling between the pseudo-label predictor and the target predictor, the inherent bias of pseudo-labels toward head classes, and artificially restricted sample weights.⁶ To mitigate this, they introduced a framework that completely decouples the adaptation process: one algorithmic pathway adapts strictly to the unlabeled data distribution to generate debiased pseudo-label predictions, while a parallel pathway adapts to the target distribution for debiased final predictions.⁶

The challenge of long-tailed distributions was approached from an entirely different vector at ICLR 2023 with the introduction of InPL (Pseudo-labeling the Inliers First).⁴¹ Recognizing the aforementioned flaw that softmax confidence is fundamentally unreliable for out-of-distribution or tail-class samples, InPL abandoned confidence-based filtering entirely. Instead, the authors leveraged energy scores—a concept rigorously developed within the out-of-distribution detection literature—to measure whether an unlabeled sample is an "inlier" relative to the current labeled training data.⁴¹ By pseudo-labeling the inliers first, regardless of their softmax confidence, InPL dramatically increased the recall for tail classes without

sacrificing overall pseudo-label precision. When integrated into existing frameworks, InPL achieved an absolute accuracy improvement of approximately 3% on CIFAR10-LT and up to 6.9% improvement over top competitors in highly imbalanced scenarios.⁴²

At ICML 2025, the Meta-Expert framework offered a highly structural solution to the long-tailed distribution mismatch problem, titled "A Square Peg in a Square Hole".³⁴ The researchers observed that in multi-expert systems, different auxiliary classifiers exhibit unique competencies depending on the data subset. For instance, a long-tailed expert network excels at predicting head classes located in high-density intervals, while a uniform expert network is vastly superior at handling samples located in medium or tail intervals.³⁴ Meta-Expert introduced a dynamic expert assignment module that actively estimates the class membership of an unlabeled sample and routes it to the most suitable expert.³⁴ A critical theoretical discovery in this research was the behavioral divergence across network depths: deep neural features tend to be highly biased toward head classes but possess strong discriminative ability, whereas shallow network features are significantly less biased but lack fine-grained discrimination.³⁴ To resolve this dichotomy, Meta-Expert introduced a Multi-depth Feature Fusion (MFF) module that blends representations from different depths, achieving state-of-the-art generalization error bounds by perfectly balancing bias mitigation and class discrimination.³⁴

Another distinct sub-problem involving multi-label data was resolved at NeurIPS 2023 with the introduction of Class-Aware Pseudo-Labeling (CAP).⁴⁶ In semi-supervised multi-label learning (SSMLL), traditional instance-aware pseudo-labeling fails catastrophically because each instance possesses an unknown number of true labels.⁴⁶ This leads to the massive introduction of false positive labels or the neglect of true positive labels.⁴⁶ CAP introduced a regularized learning framework incorporating class-distribution-aware thresholding (CAT). The researchers theoretically verified that even with a minuscule fraction of labeled examples, the empirically estimated class distribution provides a highly reliable approximation of the true target distribution.⁴⁶ CAP utilizes this insight to strictly align the predicted pseudo-label distribution with the true multi-label distribution, providing a strict generalization error bound for the multi-label space and confirming its efficacy across multiple benchmark datasets.⁴⁶

Expanding Modalities: Continuous Regression, Dense Prediction, and Graph Consistency

While the vast majority of semi-supervised learning research historically focused on discrete, categorical classification, the last four years have witnessed rigorous expansion into continuous regression outputs, dense pixel-level segmentation, and graph topological structures.¹⁰

Semi-supervised regression (SSR) poses unique mathematical and geometrical challenges. Unlike classification, continuous regression outputs lack well-defined, discrete decision boundaries that can be sharpened through standard entropy minimization.⁵² Furthermore, the

presence of heteroscedastic noise means that predictive variance and epistemic uncertainty fluctuate significantly across the continuous data domain, making naive pseudo-labeling highly susceptible to unbounded error accumulation.⁵²

At NeurIPS 2024, the RankUp architecture proposed an elegant structural solution to this geometric limitation.⁵⁶ Rather than attempting to adapt continuous outputs directly, RankUp fundamentally transformed the continuous regression task into a pairwise ranking problem by utilizing an Auxiliary Ranking Classifier (ARC) based on RankNet architectures.⁹ By comparing two unlabeled instances to determine their relative continuous order (e.g., determining which of two images depicts an older individual), RankUp allowed highly advanced, standard semi-supervised classification algorithms to be applied seamlessly to the regression space.⁹ This transformation was mathematically augmented by a Regression Distribution Alignment (RDA) technique that explicitly aligned the continuous distribution of the generated pseudo-labels.⁹ Despite its structural simplicity, RankUp achieved state-of-the-art results across a diverse array of regression benchmarks, including the BVCC audio dataset and Yelp Review text datasets, yielding major improvements in Mean Absolute Error (MAE) and Spearman's Rank Correlation Coefficient (SRCC).⁹

Researchers further addressed the specific problem of continuous variance at NeurIPS 2025 by tackling heteroscedastic noise directly.⁵² The authors proposed an uncertainty-aware pseudo-labeling framework explicitly structured as a bi-level optimization problem.⁵² The algorithm operates by jointly minimizing the empirical risk over all available labeled and pseudo-labeled data while simultaneously optimizing precise uncertainty estimates to enhance generalization specifically on the clean, labeled subset.⁵² By dynamically adjusting the gradient influence of pseudo-labels based on their calculated heteroscedastic uncertainty, this method acts as a natural dampener, preventing the runaway accumulation of errors inherent in naive continuous pseudo-labeling and demonstrating superior robustness on complex SSR benchmark datasets.⁵²

In the realm of computer vision, dense prediction, and structural topology, pseudo-labeling has been heavily adapted to handle cross-domain generalization, instance segmentation, and graph consistency. At NeurIPS 2024, researchers introduced a Connectivity-Driven Pseudo-Labeling (SeCo) framework specifically designed to fortify cross-domain segmenters.¹⁰ Real-world segmentation models often generate highly fragmented, physically impossible, and noisy pseudo-labels when deployed across significant domain shifts (e.g., synthetic to real-world environments).¹⁰ The proposed Pixel Semantic Aggregation (PSA) module systematically interacts with vision foundation models, specifically the Segment Anything Model (SAM), to aggregate disjointed pixels into continuous, geometrically logical semantic connectivities.¹⁰ This structural prior enables the model's Semantic Connectivity Correction (SCC) module to identify and correct noisy pseudo-labels at the macroscopic connectivity level rather than the microscopic pixel level, yielding massive performance improvements, such as a 13.4% increase in mIoU scores during domain adaptation from GTA5

to Cityscapes.¹⁰

Furthermore, the CAST (Contrastive Adaptation and Stillation) framework, presented at ICLR 2026, revolutionized semi-supervised knowledge distillation for the highly complex task of instance segmentation.⁶⁴ Instance segmentation demands excessively costly, per-pixel boundaries for every unique object instance.⁶⁵ CAST leverages abundant unlabeled data to compress massive, computationally expensive pre-trained Vision Foundation Models (VFMs) into highly compact, deployable expert networks.⁶⁵ The framework operates through a rigorous three-stage process: domain adaptation via contrastive calibration, knowledge transfer via multi-objective loss, and student refinement to purge residual pseudo-label bias.⁶⁵ Central to this mathematical framework is an instance-aware pixel-wise contrastive loss function that directly fuses mask and class scores.⁶⁵ This extracts highly informative negative samples and enforces strict inter-instance margins across the latent space. The results were unprecedented: a student model eleven times smaller than its VFM teacher managed to surpass the zero-shot performance of its teacher by +8.5 AP on Cityscapes and +7.1 AP on ADE20K, completely redefining the state-of-the-art for semi-supervised knowledge distillation.⁶⁵

The challenge of authentic distortion in imaging was addressed at ICML 2025 via the CPL-IQA framework.⁶⁶ Blind Image Quality Assessment (BIQA) is notoriously difficult because real-world distortions (e.g., natural motion blur, sensor noise) lack definitive ground-truth reference images.⁶⁶ CPL-IQA introduced a confidence-quantifiable pseudo-label learning method that converts Mean Opinion Score (MOS) labels to complex vector labels via entropy minimization.⁶⁶ By utilizing a manifold assumption-based label optimization strategy iteratively interleaved with model training, CPL-IQA successfully harnesses unlabeled authentically distorted images without requiring additional network complexity, standardizing semi-supervised learning for BIQA.⁶⁶

In the domain of topological data, SimMatchV2, presented at ICML 2024, extended self-training principles to graph neural networks.¹¹ By enforcing graph consistency within the pseudo-labeling loop, the model guarantees that topologically adjacent nodes within the graph structure maintain semantic consistency during the self-training propagation, effectively adapting classical semi-supervised techniques to non-Euclidean data structures.¹¹

Self-Training in the Era of Large Language Models and Foundation Models

The advent and subsequent scaling of Large Language Models (LLMs) and Vision Foundation Models (VFMs) has fundamentally altered the epistemology of self-training.⁶⁹ When the primary model generating the pseudo-labels possesses vast amounts of pre-trained, generalized world knowledge, self-training ceases to be merely a mechanism for statistical density estimation. Instead, it evolves into a mechanism of complex reasoning extraction, latent

capability discovery, and iterative autonomous scaling.

The foundational shift in this trajectory began with STaR (Self-Taught Reasoner: Bootstrapping Reasoning With Reasoning), which established that large language models can bootstrap their own complex reasoning capabilities.⁷¹ By prompting the model to generate step-by-step "chain-of-thought" rationales and subsequently filtering for reasoning chains that lead to the correct final answer, the model is iteratively fine-tuned upon its own successful logic paths.⁷¹ This creates a continuous positive feedback loop of capability enhancement, allowing models to learn from self-generated reasoning without requiring the prohibitively expensive human annotation of intermediate logical steps.⁷²

This specific capability trajectory was formalized, generalized, and scaled in the ICLR 2024 paper "Beyond Human Data: Scaling Self-Training for Problem-Solving with Language Models".⁸ The researchers demonstrated empirically that neural networks can transcend the inherent quality and diversity limitations of human-generated data, specifically in objective domains where verifiable scalar feedback is available, such as code generation and advanced mathematical problem-solving.⁸ Utilizing a self-training methodology rigorously based on the statistical Expectation-Maximization (EM) algorithm, the models autonomously generate massive quantities of synthetic reasoning paths, verify their correctness against an objective scalar environment, and optimize their parameters accordingly.⁸ This research was monumental, proving mathematically and empirically that the asymptotic performance of large language models is not strictly bounded by the volume of human data, provided an automated verification mechanism exists within the self-training loop.⁸

The concept of length generalization and complexity extrapolation was further explored at ICML 2025 in the paper "Self-Improving Transformers Overcome Easy-to-Hard and Length Generalization Challenges".⁷⁸ Standard transformer architectures notoriously struggle to generalize to sequence lengths or problem complexities unseen during their pre-training phase.⁷⁸ The researchers demonstrated that by initiating a self-training loop on simple problems and applying strict, mathematically sound filtering mechanisms—such as dynamic length thresholds and multi-path majority voting—the model could progressively synthesize and train upon increasingly complex outputs.⁷⁸ Through this purely unsupervised, autonomous curriculum, the transformer models successfully generalized from solving basic 10-digit addition problems to mastering highly complex 100-digit addition problems, as well as navigating complex mazes.⁷⁸ This proved definitively that controlled, weak-to-strong self-training curricula can exponentially expand model capabilities and overcome fundamental architectural length limitations without requiring structural modifications to the attention mechanism.⁷⁸

In the specific domain of computer vision, the NeurIPS 2025 paper "Revisiting Semi-Supervised Learning in the Era of Foundation Models" provided a definitive, highly necessary empirical study on the integration of modern VFMs (such as CLIP and DINOv2) with classical semi-supervised techniques.² The researchers discovered a counter-intuitive phenomenon:

blindly applying legacy self-training algorithms (designed for training models from scratch) to heavily pre-trained foundation models often yielded sub-optimal results, occasionally performing worse than simple supervised fine-tuning on the limited labeled subset.² However, by generating ensemble pseudo-labels simultaneously across a diverse array of distinct parameter-efficient fine-tuning (PEFT) architectures and heterogeneous foundation model backbones, they created highly robust, multi-perspective training targets.² This advanced ensemble approach actively mitigates the inherent, specific biases of any single foundation model or fine-tuning technique, providing an actionable, highly robust blueprint for how semi-supervised learning must be conducted in the modern era of massive vision models.²

The Frontiers of Epistemology: Non-Stationarity, Negative Signals, and Model Autophagy

As self-training algorithms rapidly approach the theoretical limits of pure positive reinforcement, the most recent breakthroughs from 2025 and 2026 have pivoted radically. The bleeding edge of research is now focused on leveraging training non-stationarity, actively exploiting negative degradation signals, and defining strict theoretical bounds for reinforcement learning over reasoning paths.

For over a decade, self-training operated on the assumption that high training stationarity—a condition where a model consistently predicts the exact same pseudo-label for an unlabelled sample across multiple training epochs—was a strict prerequisite for high-quality pseudo-labels.¹⁴ At ICML 2025, a paradigm-shifting paper titled "Non-Stationary Predictions May Be More Informative: Exploring Pseudo-Labels with a Two-Phase Pattern of Training Dynamics" fundamentally challenged this core tenet.¹⁴ The researchers discovered a completely new category of highly informative pseudo-labels characterized precisely by their lack of stationarity: the two-phase dynamic.¹⁴ These samples are initially predicted as one class with high confidence, but as the decision boundaries shift during training, the prediction definitively flips to another class.¹⁴ The theoretical analysis revealed that these non-stationary samples inherently exist near highly complex, deeply ambiguous decision boundaries. Consequently, they contain vastly richer class-discriminative information than easily learned, stationary samples.¹⁴ By actively isolating and incorporating these previously discarded, highly uncertain labels into the training set, the algorithm dramatically enhanced the robustness and generalization of the decision boundaries in notoriously difficult low-density regions, completely inverting traditional confidence-based selection logic.¹⁴

The looming epistemological crisis of autonomous training is the phenomenon of Model Autophagy Disorder (MAD) or Model Collapse.¹³ As generative models increasingly consume their own synthetic data—whether through self-training massive LLMs or recursive diffusion processes—they mathematically degrade in diversity, suffering from structural mode collapse and tail-class starvation.¹³ This critical failure mode was directly addressed and elegantly resolved at ICLR 2026 by the Neon framework.¹³

Neon introduced the radical, counter-intuitive concept of Negative Extrapolation from Self-Training.¹³ The researchers provided exhaustive theoretical proof demonstrating that standard mode-seeking inference samplers (which favor high-probability regions) inevitably create an exact mathematical anti-alignment between synthetic data gradients and true population data gradients.¹³ By dissecting the anisotropic dynamics of the system into an infinite hierarchy of coupled equations, they proved that self-training induces a predictable degradation vector.⁹¹ Instead of treating this degradation as a failure to be avoided, Neon actively exploits it. The base generative model deliberately undergoes a brief self-training phase purely to compute the exact directional vector of the degradation.¹³ Once computed, the algorithm algebraically extrapolates the model's neural weights backward, actively moving away from the degradation direction via a simple, computationally trivial parameter merge.¹³ This post-hoc procedure requires zero additional real data, utilizes less than 1% additional training compute, and yet achieves unprecedented state-of-the-art performance across diffusion models, flow-matching networks, and autoregressive architectures.¹³ Neon proves definitively that the noise and collapse generated by self-training can be mathematically inverted to yield pristine, highly optimized optimization signals.¹³

Finally, the vulnerability of naive self-training in complex reasoning tasks was acutely addressed at ICLR 2026 with the introduction of RESTRAIN (From Spurious Votes to Signals — Self-Training RL with Self-Penalization).⁹³ As LLMs attempt to self-train on exceedingly difficult logic puzzles, they become highly susceptible to reward hacking.⁹³ The researchers observed a catastrophic phenomenon termed "Ceiling Collapse," wherein models undergoing standard Test-Time Reinforcement Learning (TTRL) on highly complex tasks—such as the AIME24 advanced mathematics benchmark—actually began to perform significantly worse than their untrained base models.⁹³ The models were rapidly overfitting to spurious signals, discovering easy, illogical patterns that coincidentally yielded the correct final scalar answer without demonstrating genuine mathematical reasoning.⁹³

To counter this, RESTRAIN introduced an advanced self-penalization mechanism integrated directly into the reinforcement learning policy optimization objective (such as GRPO).⁹⁴ By actively identifying and explicitly penalizing spurious reasoning votes, RESTRAIN effectively turns the total absence of human ground-truth labels into highly structured, rollout-level and prompt-level negative learning signals.⁹⁴ By preventing the model from rewarding itself for correct answers achieved through faulty logic, RESTRAIN dramatically raised the performance ceiling. On the rigorous AIME25 benchmark, it boosted Pass@1 rates by over 140.7%, nearly perfectly matching the asymptotic performance of models trained exclusively on human gold-standard labels.⁹³ This establishes that for self-training to succeed in complex cognitive domains, the architecture must possess the capability to actively penalize its own epistemological blind spots.⁹⁴

Structured Inventory: Significant Research Outputs

(2022-2026)

The following table presents a structured, comprehensive compilation of the highly impactful accepted papers across NeurIPS, ICML, and ICLR that define the modern boundaries of semi-supervised self-training. It includes synthesized one-line summaries, approximate citation metrics indicating contemporary academic velocity, and a subjective evaluation of their structural impact on the discipline.

Paper Title & Venue	One-Liner Summary	Citations	Subjective Impact Evaluation
FreeMatch: Self-adaptive Thresholding for Semi-supervised Learning (ICLR 2023) ⁴	Replaces rigid static thresholds with dynamic, self-adaptive thresholding driven by the model's moving-average learning status and class fairness regularization.	~690	Transformative. FreeMatch permanently shifted the SSL landscape away from rigid heuristics like FixMatch. Its adoption of dynamic moving-average thresholds established the core baseline mechanism for modern adaptive pseudo-labeling.
SoftMatch: Addressing the Quantity-Quality Tradeoff in SSL (ICLR 2023) ⁵	Eliminates binary thresholding entirely by utilizing a continuous, truncated Gaussian weighting function to safely incorporate unconfident pseudo-labels.	~398	Highly Influential. By mathematically formalizing the quantity-quality tradeoff, SoftMatch proved that hard thresholding destroys critical gradient information. Its continuous weighting approach is fundamental for

			addressing difficult, weakly-learned classes.
Beyond Human Data: Scaling Self-Training for Problem-Solving <i>(ICLR 2024)</i> ⁸	Proves mathematically and empirically that language models can surpass human data limitations in verifiable domains using an Expectation-Maximization self-training loop.	~238	Foundational for LLMs. This paper provided the critical empirical evidence necessary to justify recursive self-improvement in LLMs. It directly paved the way for modern automated reasoning models by proving that synthetic generation scales asymptotically.
Class-Distribution-Aware Pseudo-Labeling for Semi-Supervised Multi-Label Learning <i>(NeurIPS 2023)</i> ⁴⁶	Solves the false-positive injection problem in multi-label pseudo-labeling by tightly aligning predictions with a theoretically estimated true class distribution.	~57	Definitive for Multi-Label Spaces. While restricted to the multi-label domain, this paper provided a rigorous mathematical proof that true multi-label target distributions can be reliably approximated from minuscule labeled subsets.
InPL: Pseudo-labeling the Inliers First for Imbalanced SSL <i>(ICLR 2023)</i> ⁴¹	Replaces structurally flawed softmax confidence metrics with energy-based out-of-distribution	~25	Significant Methodological Shift. By explicitly bridging out-of-distribution detection theory

	scores to preferentially select high-quality pseudo-labels for tail classes.		with semi-supervised learning, InPL highlighted the fundamental geometric flaw of softmax distributions in imbalanced data manifolds.
Bidirectional Adaptation for Robust SSL with Inconsistent Data Distributions <i>(ICML 2023)</i> ⁶	Decouples the SSL architecture into parallel pathways to independently adapt to unlabeled data distributions and target data distributions, correcting severe generalization errors.	~12	Theoretically Crucial. This paper provided the necessary mathematical explanation for why classical SSL fails in unconstrained environments. Its decoupled adaptation architecture serves as a blueprint for open-world deployments.
Connectivity-Driven Pseudo-Labeling Makes Stronger Cross-Domain Segmenters <i>(NeurIPS 2024)</i> ¹⁰	Leverages Vision Foundation Models to aggregate fragmented, noisy pixel-level predictions into coherent, continuous semantic connectivities for cross-domain adaptation.	~12	High Practical Utility. Addressing the notorious fragmentation of pseudo-labels in dense prediction, this paper successfully integrated foundation models strictly as geometric priors, massively stabilizing cross-domain

			training loops.
RankUp: Boosting Semi-Supervised Regression with an Auxiliary Ranking Classifier <i>(NeurIPS 2024)</i> ⁹	Transforms continuous scalar regression into a pairwise ranking problem via RankNet architectures, enabling advanced discrete pseudo-labeling techniques in regression spaces.	~9	Highly Innovative. Semi-supervised regression has historically lagged behind classification. By elegantly bypassing the lack of discrete decision boundaries, RankUp opened an entirely new vector for continuous SSL research.
PabLO: Improving SSL with Pseudolabeling Optimization <i>(NeurIPS 2024)</i> ⁷	Formulates confidence estimation as a strict mathematical optimization problem over functional space to automatically derive optimal scoring functions without manual tuning.	~8	Methodologically Sound. PabLO removes the final vestiges of human-designed heuristics from confidence scoring. Its strict optimization framework is highly applicable to specialized domains where standard probability distributions fail.
How Does Semi-supervised Learning with Pseudo-labelers Work? A Case Study	Provides a rigorous mathematical proof demonstrating a fundamental separation in test loss between self-trained models	~15	Theoretically Foundational. This is a landmark theoretical paper. By proving that pseudo-labeling fundamentally

(ICLR 2023)¹²	and purely supervised models under identical constraints.		alters gradient dynamics to achieve near-zero test loss where supervised models mathematically fail, it validates the entire discipline.
ASTrA: Adversarial Self-supervised Training with Adaptive-Attacks (ICLR 2025)¹⁹	Introduces an autonomous neural network that discovers optimal adversarial attack parameters through continuous exploration-exploitation during self-supervised training.	~1	Forward-Looking. Blending active adversarial robustness directly with SSL, this paper is crucial for developing models that are inherently resilient to robust overfitting while simultaneously maintaining high semi-supervised performance metrics.
Meta-Expert for Long-Tailed Semi-Supervised Learning (ICML 2025)³⁴	Dynamically routes unlabeled samples to targeted expert networks based on class estimates, utilizing multi-depth feature fusion to balance deep representation bias.	Recent	Highly Impactful for Imbalanced Data. The empirical discovery that shallow neural features resist long-tailed representation bias significantly better than deep features is profound. This will dictate how multi-expert architectures are designed.

Non-Stationary Predictions May Be More Informative <i>(ICML 2025)</i> ¹⁴	Discards the long-held requirement for training stationarity, proving mathematically that pseudo-labels exhibiting a flipping, two-phase dynamic contain superior class-discriminative data.	Recent	Paradigm Disruptor. By proving that highly uncertain, shifting predictions located exactly on the complex decision boundary are actually the most valuable pseudo-labels, this paper completely inverts traditional confidence-based selection logic.
Self-Improving Transformers Overcome Easy-to-Hard and Length Generalization <i>(ICML 2025)</i> ⁷⁸	Demonstrates that unsupervised, iterative self-training with length thresholds and majority voting allows transformers to autonomously generalize to highly complex, out-of-distribution lengths.	Recent	Critical for LLM Scaling. Length generalization is a known, fundamental failure mode of standard transformers. Proving that an unsupervised weak-to-strong curriculum solves this problem provides a highly scalable alternative to architectural overhauls.
CAST: Contrastive Adaptation and Stillation for Instance Segmentation <i>(ICLR 2026)</i> ⁶⁴	Compresses massive Vision Foundation Models into deployable, compact experts via instance-aware pixel-wise contrastive losses	Recent	SOTA for Dense Prediction. CAST successfully bridges the massive computational gap between immense foundation models and

	and continuous self-training.		edge-deployable networks, elegantly solving the extreme data scarcity problem inherent in pixel-level instance segmentation.
Neon: Negative Extrapolation From Self-Training Improves Image Generation <i>(ICLR 2026)</i> ¹³	Mathematically proves that self-training induces strict anti-alignment; autonomously extrapolates neural weights away from the self-training degradation vector to massively boost performance.	Recent	Revolutionary. Neon actively exploits the phenomenon of Model Autophagy Disorder. By using the mathematical failure of self-training as a precise negative gradient signal, it opens an entirely new sub-field of negative extrapolation in generative modeling.
RESTRAIN: Self-Training RL with Self-Penalization <i>(ICLR 2026)</i> ⁹³	Integrates aggressive self-penalization mechanisms directly into reinforcement learning to prevent models from overfitting to spurious, logically flawed patterns on complex reasoning benchmarks.	Recent	Critical for Advanced Reasoning. As models tackle increasingly difficult mathematical datasets, reward hacking becomes inevitable. RESTRAIN's ability to actively identify and penalize spurious votes is mandatory for the next generation of

			reasoning models.
Revisiting Semi-Supervised Learning in the Era of Foundation Models <i>(NeurIPS 2025)</i> ²	Evaluates legacy SSL against modern Vision Foundation Models, demonstrating that ensemble pseudo-labeling across diverse parameter-efficient fine-tuning methods actively prevents structural bias.	Recent	Highly Pragmatic. This paper prevents the SSL community from blindly applying obsolete techniques to modern foundation models. The multi-perspective ensemble approach provides the definitive baseline for how SSL should be conducted in the current era.
Semi-Supervised Regression with Heteroscedastic Pseudo-Labels <i>(NeurIPS 2025)</i> ⁵²	Formulates continuous pseudo-labeling strictly as a bi-level optimization problem to dynamically suppress the influence of highly uncertain, wildly fluctuating heteroscedastic pseudo-labels.	Recent	Mathematically Elegant. By directly modeling the continuous variance of algorithmic predictions, this paper provides a highly stable mathematical framework for regression tasks that naturally resists long-term error accumulation.

Second and Third-Order Strategic Insights and Implications

The extensive aggregation and deep analysis of this research material from NeurIPS, ICML, and ICLR reveal several profound, underlying geometric and mathematical mechanics dictating the future trajectory of machine learning. The first-order takeaway is well-established and repeatedly verified: given sufficient algorithmic safeguards and mathematical boundaries, deep neural networks are highly capable of generating their own continuous training distributions to

surpass the inherent limits of human annotation. However, the second-order insights generated from this meta-analysis reveal that the asymptotic performance of any self-training loop is strictly and invariably governed by the algorithmic handling of epistemic uncertainty and continuous distribution shift.

Historically, the entire field attempted to isolate pure signal by utilizing high-confidence thresholds. The research trajectory from 2022 to 2026 demonstrates the inherent, inescapable mathematical limits of this approach. Discarding low-confidence, noisy data starves the model of critical gradient information regarding the most complex, intertwined regions of the data manifold. SoftMatch, InPL, and the revolutionary ICML 2025 paper on Non-Stationary Predictions all independently confirm a critical, unifying second-order theme: the hardest, most uncertain, and most unstable predictions are often the most valuable data points for refining and hardening complex decision boundaries.⁵ The core optimization problem in semi-supervised learning has entirely shifted from arbitrarily *excluding* uncertainty to accurately *quantifying, modeling, and weighting* it.

Third-order insights generated specifically from the latest ICLR 2026 papers suggest a looming epistemological crisis in artificial intelligence, actively termed Model Autophagy Disorder (MAD) or Model Collapse.¹³ As generative models increasingly consume their own synthetic data loops—whether through self-training massive reasoning LLMs or recursive image diffusion models—they naturally and mathematically suffer from severe tail-class starvation and structural mode collapse.¹³ The introduction of the Neon framework and the RESTRAIN architecture reveals the necessary, permanent solution: the active utilization of negative, adversarial signals.¹³ Because standard mode-seeking samplers mathematically induce an absolute anti-alignment between synthetic gradients and true underlying population gradients, infinite positive self-training is mathematically guaranteed to degrade the neural network eventually.⁸⁸

To achieve continuous, unbounded self-improvement, models must now be structurally programmed to recognize, map, and penalize their own failure modes. Neon achieves this by actively computing the degradation vector and extrapolating neural weights aggressively away from the self-trained degradation.¹³ RESTRAIN achieves this by exhaustively mapping and explicitly penalizing spurious reasoning paths that yield correct final answers through demonstrably faulty, hallucinated logic.⁹³ Therefore, the bleeding edge frontier of self-training is no longer about maximizing confidence in generated pseudo-labels; it is entirely about actively identifying, modeling, and mathematically penalizing the model's own epistemological blind spots. Furthermore, the boundaries between traditional semi-supervised learning and continuous reinforcement learning are rapidly dissolving. As demonstrated by STaR and RESTRAIN, pseudo-labels in the era of massive foundation models are no longer static, categorical target vectors; they are continuous, dynamic reward signals actively shaping long-horizon, complex reasoning trajectories across infinite latent spaces.

Works cited

1. Self-Training: A Survey - arXiv, accessed April 23, 2026, <https://arxiv.org/html/2202.12040v6>
2. Revisiting Semi-Supervised Learning in the Era of Foundation Models - OpenReview, accessed April 23, 2026, [https://openreview.net/forum?id=yDh9s1qPzB&referrer=%5Bthe%20profile%20of%20Zheda%20Mai%5D\(%2Fprofile%3Fid%3D~Zheda_Mai1\)](https://openreview.net/forum?id=yDh9s1qPzB&referrer=%5Bthe%20profile%20of%20Zheda%20Mai%5D(%2Fprofile%3Fid%3D~Zheda_Mai1))
3. ICML 2024 Papers, accessed April 23, 2026, <https://icml.cc/virtual/2024/papers.html>
4. (PDF) FreeMatch: Self-adaptive Thresholding for Semi-supervised Learning - ResearchGate, accessed April 23, 2026, https://www.researchgate.net/publication/360642554_FreeMatch_Self-adaptive_Thresholding_for_Semi-supervised_Learning
5. SOFTMATCH: ADDRESSING THE QUANTITY-QUALITY TRADE-OFF IN SEMI-SUPERVISED LEARNING - OpenReview, accessed April 23, 2026, <https://openreview.net/pdf?id=yMt1zQXBDiF>
6. ICML 2023 Bidirectional Adaptation for Robust Semi-Supervised Learning with Inconsistent Data Distributions Oral, accessed April 23, 2026, <https://icml.cc/virtual/2023/oral/25533>
7. PabLO: Improving Semi-Supervised Learning with Pseudolabeling Optimization, accessed April 23, 2026, <https://neurips.cc/virtual/2024/104518>
8. Beyond Human Data: Scaling Self-Training for Problem-Solving with Language Models, accessed April 23, 2026, <https://iclr.cc/virtual/2024/22448>
9. RankUp: Boosting Semi-Supervised Regression with an Auxiliary Ranking Classifier, accessed April 23, 2026, <https://openreview.net/forum?id=d2lPM1Aczs>
10. Connectivity-Driven Pseudo-Labeling Makes Stronger Cross-Domain Segmenters - NIPS papers, accessed April 23, 2026, https://proceedings.neurips.cc/paper_files/paper/2024/file/8fd1a81c882cd45f64958da6284f4a3f-Paper-Conference.pdf
11. Feature Space Renormalization for Semi-supervised Learning - ResearchGate, accessed April 23, 2026, https://www.researchgate.net/publication/402744807_Feature_Space_Renormalization_for_Semi-supervised_Learning
12. How Does Semi-supervised Learning with Pseudo-labelers Work? A Case Study, accessed April 23, 2026, <https://iclr.cc/virtual/2023/poster/10909>
13. Neon: Negative Extrapolation From Self-Training Improves Image Generation - OpenReview, accessed April 23, 2026, <https://openreview.net/forum?id=kpLRYtPGt3>
14. ICML Poster Non-Stationary Predictions May Be More Informative: Exploring Pseudo-Labels with a Two-Phase Pattern of Training Dynamics, accessed April 23, 2026, <https://icml.cc/virtual/2025/poster/44953>
15. How Does Semi-supervised Learning with Pseudo-labelers Work? A ..., accessed April 23, 2026, <https://openreview.net/forum?id=Dzmd-Cc8OI>
16. HOW DOES SEMI-SUPERVISED LEARNING WITH PSEUDO-LABELERS WORK? A CASE STUDY - OpenReview, accessed April 23, 2026, <https://openreview.net/pdf?id=Dzmd-Cc8OI>

17. How Does Semi-supervised Learning with Pseudo-labelers Work? A Case Study, accessed April 23, 2026, <https://www.semanticscholar.org/paper/How-Does-Semi-supervised-Learning-with-Work-A-Case-Kou-Chen/3ec356c38f0b2a0828e475c8d73e907691251b6f>
18. NeurIPS Poster Pseudo-Labeling for Kernel Ridge Regression under Covariate Shift, accessed April 23, 2026, <https://neurips.cc/virtual/2025/poster/122007>
19. Automated Discovery of Adaptive Attacks on Adversarial Defenses - Semantic Scholar, accessed April 23, 2026, <https://www.semanticscholar.org/paper/Automated-Discovery-of-Adaptive-Attacks-on-Defenses-Yao-Bielik/9474a3056dbae6952218997fd8189622a68bec0d>
20. Mubarak Shah - Center for Research in Computer Vision - University of Central Florida, accessed April 23, 2026, <https://www.crcv.ucf.edu/person/mubarak-shah/>
21. Prakash Chandra Chhipa, accessed April 23, 2026, <https://prakashchhipa.github.io/>
22. SoftMatch: Addressing the Quantity-Quality Tradeoff in Semi-supervised Learning, accessed April 23, 2026, <https://openreview.net/forum?id=ymt1zQXBDiF>
23. PabLO: Improving Semi-Supervised Learning with Pseudolabeling Optimization - 玻尔, accessed April 23, 2026, <https://www.bohrium.com/paper-details/pablo-improving-semi-supervised-learning-with-pseudolabeling-optimization/1134739577047613440-2000000>
24. FreeMatch: Self-adaptive Thresholding for Semi-supervised Learning | OpenReview, accessed April 23, 2026, https://openreview.net/forum?id=PDrUPTXJI_A
25. ReliaMatch: Semi-Supervised Classification with Reliable Match - MDPI, accessed April 23, 2026, <https://www.mdpi.com/2076-3417/13/15/8856>
26. FreeMatch: Self-adaptive Thresholding for Semi-supervised Learning - arXiv, accessed April 23, 2026, <https://arxiv.org/pdf/2205.07246>
27. Hao Chen - Google Scholar, accessed April 23, 2026, <https://scholar.google.com/citations?user=tktgkhwAAAAJ&hl=en>
28. [2301.10921] SoftMatch: Addressing the Quantity-Quality Trade-off in Semi-supervised Learning - arXiv, accessed April 23, 2026, <https://arxiv.org/abs/2301.10921>
29. Marios Savvides - Google Scholar, accessed April 23, 2026, https://scholar.google.com/citations?user=mWpu_ooAAAAJ&hl=en
30. Yue Fan - Google Scholar, accessed April 23, 2026, <https://scholar.google.com/citations?user=TNMSbOkAAAAJ&hl=en>
31. PabLO: Improving Semi-Supervised Learning with Pseudolabeling Optimization - Harit Vishwakarma, accessed April 23, 2026, <https://harit7.github.io/assets/pdf/pablo-icml-dmlr-ws-2024.pdf>
32. Yi Chen - Google 学术搜索, accessed April 23, 2026, <https://scholar.google.co.id/citations?user=QoO6pMEAAAAJ&hl=zh-CN>
33. Yi Chen - Google Scholar, accessed April 23, 2026, <https://scholar.google.com/citations?user=QoO6pMEAAAAJ&hl=en>
34. ICML Poster A Square Peg in a Square Hole: Meta-Expert for Long-Tailed Semi-Supervised Learning - ICML 2026, accessed April 23, 2026,

- <https://icml.cc/virtual/2025/poster/44441>
35. A Square Peg in a Square Hole: Meta-Expert for Long-Tailed Semi-Supervised Learning | OpenReview, accessed April 23, 2026, <https://openreview.net/forum?id=h0ZeiDRN8A>
 36. A Square Peg in a Square Hole: Meta-Expert for Long-Tailed Semi-Supervised Learning - GitHub, accessed April 23, 2026, <https://raw.githubusercontent.com/mlresearch/v267/main/assets/hou25d/hou25d.pdf>
 37. Oral B5 Self/Semi-Supervised Learning and Interpretability / Observing Aspects of NN, accessed April 23, 2026, <https://icml.cc/virtual/2023/session/25594>
 38. Bidirectional Adaptation for Robust Semi-Supervised Learning with Inconsistent Data Distributions, accessed April 23, 2026, <https://proceedings.mlr.press/v202/jia23a.html>
 39. Lin-Han Jia - Google Scholar, accessed April 23, 2026, <https://scholar.google.com/citations?user=2Z2bdDcAAAAJ&hl=en>
 40. Jie-Jing Shao - Google Scholar, accessed April 23, 2026, <https://scholar.google.com.hk/citations?user=k1tEDpQAAAAJ&hl=en>
 41. InPL: Pseudo-labeling the Inliers First for Imbalanced Semi-supervised Learning, accessed April 23, 2026, <https://iclr.cc/virtual/2023/poster/11998>
 42. InPL: Pseudo-labeling the Inliers First for Imbalanced Semi-supervised Learning, accessed April 23, 2026, https://www.researchgate.net/publication/369198917_InPL_Pseudo-labeling_the_inliers_First_for_Imbalanced_Semi-supervised_Learning
 43. InPL: Pseudo-labeling the Inliers First for Imbalanced Semi-supervised Learning - Liner, accessed April 23, 2026, <https://liner.com/review/inpl-pseudolabeling-the-inliers-first-for-imbalanced-semisupervised-learning>
 44. Zhuoran Yu - Google Scholar, accessed April 23, 2026, <https://scholar.google.com/citations?user=txxhxREAAAAJ&hl=en>
 45. A Square Peg in a Square Hole: Meta-Expert for Long-Tailed Semi-Supervised Learning, accessed April 23, 2026, <https://arxiv.org/html/2505.16341v1>
 46. Class-Distribution-Aware Pseudo-Labeling for Semi-Supervised Multi-Label Learning, accessed April 23, 2026, <https://neurips.cc/virtual/2023/poster/71510>
 47. Class-Distribution-Aware Pseudo-Labeling for Semi-Supervised Multi-Label Learning | Request PDF - ResearchGate, accessed April 23, 2026, https://www.researchgate.net/publication/402852361_Class-Distribution-Aware_Pseudo-Labeling_for_Semi-Supervised_Multi-Label_Learning
 48. Class-Distribution-Aware Pseudo-Labeling for Semi-Supervised Multi-Label Learning | Read Paper on Bytez, accessed April 23, 2026, https://bytez.com/docs/neurips/71510/paper?_c=eyJ2IjoxLCJyZWxhdGVkljpbImNvZGUlLCJyZWZlcmVuY2VzIiwiaY29uZmVyeZW5jZSIdfQ%3D%3D
 49. Track: Poster Session 4 - NeurIPS 2026, accessed April 23, 2026, <https://neurips.cc/virtual/2023/session/74072>
 50. Ming-Kun Xie - Google 学术搜索, accessed April 23, 2026, <https://scholar.google.com/citations?user=DxUxwnMAAAAJ&hl=zh-CN>

51. Jia-Hao Xiao - Google 學術搜尋, accessed April 23, 2026, https://scholar.google.com/citations?user=p_4b2l8AAAAJ&hl=zh-TW
52. NeurIPS Poster Semi-Supervised Regression with Heteroscedastic Pseudo-Labels, accessed April 23, 2026, <https://neurips.cc/virtual/2025/poster/118197>
53. Chang Xu 0002 - DBLP, accessed April 23, 2026, <https://dblp.org/pid/97/2966-2>
54. Semi-Supervised Regression with Heteroscedastic Pseudo-Labels, accessed April 23, 2026, <https://neurips.cc/media/neurips-2025/Slides/118197.pdf>
55. Semi-Supervised Regression with Heteroscedastic Pseudo-Labels - arXiv, accessed April 23, 2026, <https://arxiv.org/html/2510.15266v1>
56. RankUp: Towards High-rank Representations for Large Scale Advertising Recommender Systems - arXiv, accessed April 23, 2026, <https://arxiv.org/html/2604.17878v1>
57. RankUp: Boosting Semi-Supervised Regression with an Auxiliary Ranking Classifier - NIPS papers, accessed April 23, 2026, https://proceedings.neurips.cc/paper_files/paper/2024/file/c26a8494fe31695db965ae8b7244b7c1-Paper-Conference.pdf
58. RankUp: Boosting Semi-Supervised Regression with an Auxiliary Ranking Classifier - arXiv, accessed April 23, 2026, <https://arxiv.org/html/2410.22124v1>
59. Pin-Yen Huang - Google Scholar, accessed April 23, 2026, <https://scholar.google.com.tw/citations?user=nQdph2MAAAAJ>
60. Connectivity-Driven Pseudo-Labeling Makes Stronger Cross-Domain Segmenters, accessed April 23, 2026, <https://www.proceedings.com/079017-2506.html>
61. Connectivity-Driven Pseudo-Labeling Makes Stronger Cross-Domain Segmenters | Request PDF - ResearchGate, accessed April 23, 2026, https://www.researchgate.net/publication/397200680_Connectivity-Driven_Pseudo-Labeling_Makes_Stronger_Cross-Domain_Segmenters
62. SeCoV2: Semantic Connectivity-Driven Pseudo-Labeling for Robust Cross-Domain Semantic Segmentation - IEEE Computer Society, accessed April 23, 2026, <https://www.computer.org/csdl/journal/tp/2025/11/11120362/2917jgEcOk0>
63. dong zhao - Google 学术搜索, accessed April 23, 2026, https://scholar.google.com/citations?user=j_UjUUUAAAAJ&hl=zh-CN
64. Better (pseudo-)labels for semi-supervised instance segmentation - ICLR 2026, accessed April 23, 2026, <https://iclr.cc/virtual/2024/21914>
65. CAST: Contrastive Adaptation and Distillation for Semi-Supervised ..., accessed April 23, 2026, <https://openreview.net/forum?id=jqLqBG6d2q>
66. Semi-Supervised Blind Quality Assessment with Confidence-quantifiable Pseudo-label Learning for Authentic Images - ICML 2026, accessed April 23, 2026, <https://icml.cc/virtual/2025/poster/46701>
67. GBVSSL: Contrastive Semi-Supervised Learning Based on Generalized Bias-Variance Decomposition - ResearchGate, accessed April 23, 2026, https://www.researchgate.net/publication/381355530_GBVSSL_Contrastive_Semi-Supervised_Learning_Based_on_Generalized_Bias-Variance_Decomposition

68. InterLUDE: Interactions Between Labeled And Unlabeled Data To Enhance Semi-Supervised Learning - Scholars' Mine, accessed April 23, 2026, https://scholarsmine.mst.edu/cgi/viewcontent.cgi?article=3071&context=comsci_facwork
69. About Me - Zheda Mai's Homepage, accessed April 23, 2026, <https://zheda-mai.github.io/>
70. Lechao Xiao - Google Scholar, accessed April 23, 2026, <https://scholar.google.com/citations?user=fvwzUnIAAAAJ&hl=en>
71. Compositional Neuro-Symbolic Reasoning - arXiv, accessed April 23, 2026, <https://arxiv.org/html/2604.02434v1>
72. luo-junyu/Awesome-Agent-Papers - GitHub, accessed April 23, 2026, <https://github.com/luo-junyu/awesome-agent-papers>
73. Self-Taught Reasoner: Bootstrapping Logic | PDF | Applied Mathematics - Scribd, accessed April 23, 2026, <https://www.scribd.com/document/753313852/2203-14465v2>
74. Through the Lens of Contrast: Self-Improving Visual Reasoning in VLMs | OpenReview, accessed April 23, 2026, <https://openreview.net/forum?id=ZymCPON45y>
75. Ankesh Anand - Google Scholar, accessed April 23, 2026, <https://scholar.google.com/citations?user=HEozmkMAAAAJ&hl=en>
76. How to Generate and Use Synthetic Data for Finetuning - Eugene Yan, accessed April 23, 2026, <https://eugeneyan.com/writing/synthetic/>
77. Yamini Bansal - Google Scholar, accessed April 23, 2026, <https://scholar.google.com/citations?user=uj1OijkAAAAJ&hl=en>
78. ICML Poster Self-Improving Transformers Overcome Easy-to-Hard and Length Generalization Challenges, accessed April 23, 2026, <https://icml.cc/virtual/2025/poster/44828>
79. Self-Supervised Transformers as Iterative Solution Improvers for Constraint Satisfaction, accessed April 23, 2026, <https://icml.cc/virtual/2025/poster/45737>
80. Proceedings of Machine Learning Research | Proceedings of the 42nd International Conference on Machine Learning Held in Vancouver Convention Center, Vancouver, Canada on 13-19 July 2025 Published as Volume 267 by the Proceedings of Machine Learning Research on 06 October 2025. Volume Edited by: Aarti Singh Maryam Fazel Daniel Hsu Simon Lacoste-Julien Felix Berkenkamp Tegan, accessed April 23, 2026, <https://proceedings.mlr.press/v267/>
81. Can Large Reasoning Models Self-Train? - arXiv, accessed April 23, 2026, <https://arxiv.org/html/2505.21444v2>
82. Self-Improving Transformers Overcome Easy-to-Hard and Length Generalization Challenges | OpenReview, accessed April 23, 2026, <https://openreview.net/forum?id=ZtXOMBT6mf>
83. NeurIPS 2025 Papers, accessed April 23, 2026, <https://neurips.cc/virtual/2025/papers.html>
84. NeurIPS Poster Revisiting Semi-Supervised Learning in the Era of Foundation Models, accessed April 23, 2026, <https://neurips.cc/virtual/2025/poster/115173>
85. Non-Stationary Predictions May Be More Informative: Exploring Pseudo-Labels

- with a Two-Phase Pattern of Training Dynamics - GitHub, accessed April 23, 2026, <https://raw.githubusercontent.com/mlresearch/v267/main/assets/pei25a/pei25a.pdf>
86. Downloads 2025 - ICML 2026, accessed April 23, 2026, <https://icml.cc/Downloads/2025>
 87. Non-Stationary Predictions May Be More Informative: Exploring Pseudo-Labels with a Two-Phase Pattern of Training Dynamics | OpenReview, accessed April 23, 2026, <https://openreview.net/forum?id=Xcnbtxns0V>
 88. Neon: Negative Extrapolation From Self-Training Improves Image Generation - ICLR 2026, accessed April 23, 2026, <https://iclr.cc/virtual/2026/oral/10007747>
 89. ICLR 2026 Thursday 04/23, accessed April 23, 2026, <https://iclr.cc/virtual/2026/day/4/23>
 90. Downloads - ICLR 2026, accessed April 23, 2026, <https://iclr.cc/Downloads/2026>
 91. Orals - ICLR 2026, accessed April 23, 2026, <https://iclr.cc/virtual/2026/events/oral>
 92. VITA-Group repositories - GitHub, accessed April 23, 2026, <https://github.com/orgs/VITA-Group/repositories>
 93. RESTRAIN: From Spurious Votes to Signals — Self-Training RL with Self-Penalization, accessed April 23, 2026, <https://openreview.net/forum?id=87ySF7viys>
 94. restrain: from spurious votes to signals — self-driven rl with self-penalization - OpenReview, accessed April 23, 2026, <https://openreview.net/pdf?id=87ySF7viys>
 95. Understanding and Mitigating Spurious Signal Amplification in Test-Time Reinforcement Learning for Math Reasoning - arXiv, accessed April 23, 2026, <https://arxiv.org/html/2604.21327v1>
 96. RESTRAIN: From Spurious Votes to Signals — Self-Training RL with, accessed April 23, 2026, <https://liner.com/review/restrain-from-spurious-votes-to-signals-selftraining-rl-with-selfpenalization>
 97. ICLR 2026 Papers, accessed April 23, 2026, <https://iclr.cc/virtual/2026/papers.html>
 98. Yiwen Kou - Google Scholar, accessed April 23, 2026, <https://scholar.google.com/citations?user=BSQAMzkAAAJ&hl=en>